

SHIVA KUMAR GAUTAM

+91 9369637611 • kumarshiva0707@gmail.com • linkedin.com/in/shivakumar07 • github.com/shivakr07

SUMMARY

Data Engineer with 2 years of experience designing enterprise data platforms, developing scalable data pipelines, and enabling data-driven decision making. Skilled in data modeling, analytics engineering, performance optimization, and AI-powered solutions, with experience building reliable and scalable systems for large-scale business workloads.

TECHNICAL SKILLS

Programming: Python, C#, SQL, Java

Data Engineering: Microsoft Fabric, Azure Data Factory (ADF), Azure Synapse Analytics, PySpark, Apache Spark, Apache Kafka, Delta Lake, Kusto (ADX/KQL)

Analytics & AI: Power BI, Semantic Models, DAX, Azure OpenAI, Copilot Studio, Fabric AI Agents, RAG

Cloud & Tools: Azure, Azure SQL DB, Azure Data Lake Storage (ADLS Gen2), Azure Functions, ADO, Git

PROFESSIONAL EXPERIENCE

MAQ Software, Hyderabad — Software Engineer

Jan 2025 – Present

- Developed **Athena**, an enterprise **RAG-based AI Copilot** platform serving **13,000+ Microsoft employees**, enabling real-time insights across Azure capacity domains while eliminating manual Kusto querying. Built a parent agent in **Copilot Studio** orchestrating **5 Fabric domain-specific AI agents** over **Kusto databases, Fabric Lakehouse, and Semantic Models**, with automatic **KQL generation** for response validation [**Received Appreciation**] ↗..
- Developed **metadata-driven ETL pipelines** using **Microsoft Fabric, Azure Data Factory, Synapse, and PySpark**, implementing **Medallion Architecture (Bronze, Silver, Gold)** to build reusable, configuration-driven data processing workflows for Azure capacity analytics.
- Implemented a reusable **Kusto cluster scaling framework** in **Synapse notebooks**, combining adaptive recursive time-window downshifting with dynamic vertical and horizontal cluster scaling to reliably process high-volume telemetry datasets while reducing compute costs by **~30%**.
- Built centralized **Fabric Lakehouse** datasets and **Semantic Models** using **Star** and **Snowflake** schema design, enabling governed enterprise reporting and self-service analytics across multiple Power BI dashboards.
- Improved reporting performance by optimizing **DAX measures**, Semantic Models, and **partition-based incremental refresh**, reducing dataset refresh time by **55%** while improving dashboard responsiveness.
- Developed **Operational Dashboards** monitoring **300+ Azure Data Factory pipelines**, providing visibility into pipeline health, data freshness, SLA compliance, failures, and production operations.

HashedIn by Deloitte, Bengaluru — Software Engineer Trainee

Oct 2024 – Jan 2025

- Developed batch and streaming **ETL pipelines** using **Apache Kafka, PySpark Structured Streaming, and Delta Lake**, implementing **Medallion Architecture (Bronze, Silver, Gold)** to process data for analytical workloads.
- Built scalable data transformation pipelines using **PySpark**, performing schema validation, data cleansing, aggregations, and business-rule transformations to improve data quality and support downstream analytics.
- Developed a **multi-agent SDLC automation system** using **LangChain, LLMs, and RAG**, enabling AI agents to assist across software development lifecycle stages including requirements analysis, code generation, testing, and documentation.

EDUCATION

B.Tech in Computer Science and Artificial Intelligence

Graduated: June 2024

Dr. A.P.J. Abdul Kalam Technical University

8.5 CGPA

PSIT College of Engineering, Kanpur, Uttar Pradesh

CERTIFICATIONS

DP-700 – Fabric Data Engineer Associate ↗

DP-600 – Fabric Analytics Engineer Associate ↗

DP-800 – Querying Data with Microsoft Transact-SQL ↗
PostgreSQL ↗

AB620 – Build AI Apps with Azure Database for